

Survival Analysis of a Randomised Cancer Trial

Does adjuvant chemotherapy delay recurrence and prolong survival in colon cancer?

Kingsley Amegah

June 2026

Contents

1	Background and question	1
2	The data and randomisation check	2
3	Kaplan-Meier survival curves	3
3.1	Overall survival (time to death)	3
3.2	Recurrence-free survival (time to recurrence)	4
4	Cox proportional-hazards model	5
5	Checking the proportional-hazards assumption	7
6	A machine-learning benchmark: Random Survival Forest	8
7	Results	10
8	Conclusions and clinical implications	10
9	Limitations	10

1 Background and question

A central job of data science in health is to answer one question rigorously: **does a treatment actually work, and for whom?** The cleanest evidence comes from a randomised controlled trial (RCT), where patients are assigned to treatments by chance so the groups are comparable.

This report analyses a real, published RCT: the **colon cancer adjuvant-therapy trial** (Moertel and colleagues). After surgery, **929 patients** with stage B or C colon cancer were randomly assigned to one of three arms:

- **Obs**, observation only (the control),
- **Lev**, levamisole,

- **Lev+5FU**, levamisole plus fluorouracil.

Two outcomes were followed over time: **recurrence** of the cancer and **death**. The analytical challenge is that many patients had not had the event by the end of follow-up (their times are *censored*), so ordinary averages do not work. This is what **survival analysis** is built for.

2 The data and randomisation check

Because treatment was randomised, the arms should look similar at baseline. A balance table is the first thing a trial statistician checks.

```
bt <- os |>
  select(rx, age, sex, nodes, node4, differ, extent, obstruct, perfor) |>
  tbl_summary(by = rx,
    label = list(age ~ "Age (years)", nodes ~ "Positive nodes",
      node4 ~ "Nodal burden", differ ~ "Differentiation",
      extent ~ "Local extent", obstruct ~ "Obstruction",
      perfor ~ "Perforation", sex ~ "Sex")) |>
  add_p() |>
  modify_header(label ~ "**Characteristic**") |>
  modify_caption("**Table 1. Baseline characteristics by treatment arm**")
cat(strip_dash(as_kable(bt)), sep = "\n")
```

Table 1: **Table . Baseline characteristics by treatment arm**

Characteristic	Obs N =	Lev N =	Lev+FU N =	p-vale
Age (years)	6 (, 68)	6 (, 69)	6 (, 7)	.8
Sex				.8
Female	9 (7%)	(%)	6 (%)	
Male	66 (%)	77 (7%)	(6%)	
Positive nodes	(,)	(,)	(,)	.6
Unknown		6	9	
Nodal brden				.7
- nodes	8 (7%)	(7%)	(7%)	
> nodes	87 (8%)	89 (9%)	79 (6%)	
Differentiation				.
Moderate	9 (7%)	9 (7%)	(7%)	
Poor	(7%)	(%)	(8%)	
Well	7 (8.8%)	7 (%)	9 (9.7%)	
Unknown	7		6	
Local extent				.
Contigos	(6.%)	(.9%)	(.6%)	
Mscle	8 (%)	6 (%)	(%)	
Serosa	9 (79%)	9 (8%)	(8%)	
Sbmcosa	8 (.%)	(.%)	(.%)	
Obstrction	6 (%)	6 (%)	(8%)	.7
Perforation	9 (.9%)	(.%)	8 (.6%)	>.9


The arms are well balanced: no baseline characteristic differs meaningfully across them (all p-values are large). That is exactly what randomisation should deliver, and it means any later difference in outcomes can be attributed to treatment rather than to sicker patients landing in one arm.

3 Kaplan-Meier survival curves

The Kaplan-Meier estimator turns censored follow-up times into a survival curve: the share of patients still event-free as time passes. We compare arms with the **log-rank test**.

3.1 Overall survival (time to death)

```
fit_os <- survfit(Surv(time_yrs, status) ~ rx, data = os)
ggsurvplot(fit_os, data = os, palette = PAL, conf.int = TRUE,
  risk.table = TRUE, pval = TRUE, censor = FALSE,
  xlab = "Years from randomisation", ylab = "Overall survival",
  legend.title = "Arm", legend.labs = c("Obs", "Lev", "Lev+5FU"),
  risk.table.height = 0.28, ggtheme = theme_minimal(base_size = 12))
```



```
lr_os <- survdiff(Surv(time_yrs, status) ~ rx, data = os)
p_os <- 1 - pchisq(lr_os$chisq, df = length(lr_os$n) - 1)
med <- summary(fit_os)$table[, "median"]
cat("Log-rank p-value (overall survival):", signif(p_os, 3), "\n")
```

```
## Log-rank p-value (overall survival): 0.0029
```

```
cat("Median survival not reached where shown as NA (over half still alive).\n")
```

```
## Median survival not reached where shown as NA (over half still alive).
```


```
print(round(med, 2))
```

```
##      rx=Obs      rx=Lev rx=Lev+5FU
##      5.70      5.89      NA
```

The **Lev+5FU** curve sits clearly above the other two, and the log-rank test confirms the arms differ significantly. More than half of the Lev+5FU patients were still alive at the end of follow-up, so their median survival is not even reached. Levamisole alone tracks the control arm closely.

3.2 Recurrence-free survival (time to recurrence)

```
fit_rfs <- survfit(Surv(time_yrs, status) ~ rx, data = rfs)
ggsurvplot(fit_rfs, data = rfs, palette = PAL, conf.int = TRUE,
  risk.table = TRUE, pval = TRUE, censor = FALSE,
  xlab = "Years from randomisation", ylab = "Recurrence-free survival",
  legend.title = "Arm", legend.labs = c("Obs", "Lev", "Lev+5FU"),
  risk.table.height = 0.28, ggtheme = theme_minimal(base_size = 12))
```



The same story appears earlier and even more strongly for recurrence: the combination arm keeps far more patients cancer-free, and most of its benefit is in preventing the cancer from coming back in the first place.

4 Cox proportional-hazards model

The Kaplan-Meier curves compare arms on their own. The **Cox model** estimates the effect of treatment while adjusting for prognostic factors, and reports a **hazard ratio (HR)**: the relative rate of death. HR below 1 means lower risk.

```
cox <- coxph(Surv(time_yrs, status) ~ rx + age + sex + nodes + extent +
             differ + obstruct + perfor, data = os)
ct <- tbl_regression(cox, exponentiate = TRUE,
                    add_estimate_to_reference_rows = TRUE,
                    label = list(rx ~ "Treatment arm", age ~ "Age",
                                sex ~ "Sex", nodes ~ "Positive nodes",
                                extent ~ "Local extent", differ ~ "Differentiation",
                                obstruct ~ "Obstruction", perfor ~ "Perforation")) |>
```

```
modify_caption("**Table 2. Adjusted hazard ratios for death (multivariable Cox model)**")
cat(strip_dash(as_kable(ct)), sep = "\n")
```

Table 2: Table . Adjsted hazard ratios for death (mltivari-
able Cox model)


Characteristic	HR	9% CI	p-vale
Treatment arm			
Obs	.		
Lev	.9	.7, .	.
Lev+FU	.67	., .86	.
Age	.	., .	.9
Sex			
Female	.		
Male	.98	.8, .8	.8
Positive nodes	.9	.7, .	<.
Local extent			
Contigos	.		
Mscle	.	., .7	.
Serosa	.7	.6, .7	.
Sbmcosa	.	.6, .7	.
Differentiation			
Moderate	.		
Poor	.	., .7	.8
Well	.9	.79, .	.6
Obstrction			
No	.		
Yes	.6	., .6	.
Perforation			
No	.		
Yes	.	.6, .77	.9

```
hr <- tidy(cox, exponentiate = TRUE, conf.int = TRUE) |>
  filter(term == "rxLev+5FU")
cat(sprintf("Lev+5FU vs Obs: HR = %.2f (95%% CI %.2f to %.2f), p = %.4f\n",
  hr$estimate, hr$conf.low, hr$conf.high, hr$p.value))
```

```
## Lev+5FU vs Obs: HR = 0.67 (95% CI 0.53 to 0.86), p = 0.0012
```

After adjusting for age, nodal burden, tumour extent and differentiation, the combination arm reduces the rate of death by roughly **one third** compared with observation, and the effect is statistically significant. The number of positive lymph nodes and the local extent of the tumour are the strongest prognostic drivers, as clinical knowledge would predict.

```
ggforest(cox, data = os, main = "Hazard ratios for death (95% CI)")
```



5 Checking the proportional-hazards assumption


The Cox model assumes a treatment's effect on the hazard is roughly constant over time. We test this with scaled Schoenfeld residuals.

```
zph <- cox.zph(cox)
print(zph)
```

```
##          chisq df      p
## rx          2.9209  2 0.23213
## age         1.1807  1 0.27722
## sex         0.3537  1 0.55202
## nodes       0.0203  1 0.88662
## extent      7.1654  3 0.06681
## differ     16.8839  2 0.00022
```

```
## obstruct 6.8003 1 0.00911
## perfor 0.5598 1 0.45434
## GLOBAL 37.9215 12 0.00016
```

```
plot(zph[1], col = ACC, lwd = 2)
abline(h = 0, lty = 2, col = "grey50")
title("Schoenfeld residuals for treatment arm")
```



The global test and the treatment term are not significant, so the proportional-hazards assumption is reasonable for the treatment effect: the hazard ratios above are a fair summary across the follow-up period.

6 A machine-learning benchmark: Random Survival Forest

Cox is the clinical standard, but it assumes effects are linear and additive. A **Random Survival Forest (RSF)** is a machine-learning model that captures non-linearities and interactions automatically. We compare the two on discrimination, measured by the concordance index (C-index, higher is better), and read off which factors the forest finds most important.

```
set.seed(42)
rsf_dat <- os |>
  select(time_yrs, status, rx, age, sex, nodes, extent, differ,
         obstruct, perfor) |> na.omit()
rsf <- rfsrc(Surv(time_yrs, status) ~ ., data = rsf_dat,
            ntree = 600, importance = TRUE)
```




```
c_rsf <- 1 - rsf$err.rate[rsf$ntree]
c_cox <- summary(cox)$concordance[1]
cat(sprintf("C-index, Cox model           : %.3f\n", c_cox))
```

```
## C-index, Cox model           : 0.663
```

```
cat(sprintf("C-index, Random Survival Forest: %.3f\n", c_rsf))
```

```
## C-index, Random Survival Forest: 0.621
```

```
vi <- sort(rsf$importance, decreasing = TRUE)
vidf <- data.frame(var = names(vi), imp = as.numeric(vi))
ggplot(vidf, aes(reorder(var, imp), imp)) +
  geom_col(fill = "teal") + coord_flip() +
  labs(title = "Random Survival Forest: variable importance",
       x = NULL, y = "Permutation importance (higher = more predictive)")
```



The two models reach **very similar discrimination**, which is reassuring: the straightforward Cox model is not leaving much on the table here. The forest ranks the same drivers at the top (nodal burden, tumour extent, age), and it also picks up treatment arm as informative. When a simple, interpretable model matches a flexible one, the simple model is usually the right choice to report.

7 Results

This section reports the findings in the style of a clinical journal.

Of 929 randomised patients, 452 deaths and 468 recurrences were observed over a maximum follow-up of approximately nine years. Baseline characteristics were well balanced across the three arms (Table 1), as expected under randomisation, so the between-arm comparisons are unlikely to be confounded by measured prognostic factors.

Overall survival differed significantly between arms (log-rank $P = 0.003$). Median overall survival was 5.7 years in the observation arm and was not reached in the levamisole-plus-fluorouracil arm (that is, more than half of those patients were still alive at the end of follow-up). In the multivariable Cox proportional-hazards model (Table 2), levamisole plus fluorouracil was associated with a significantly lower hazard of death than observation (adjusted hazard ratio [HR] 0.67; 95% confidence interval [CI], 0.53 to 0.86; $P = 0.001$), corresponding to an approximately 33% relative reduction in the rate of death. Levamisole alone conferred no significant benefit (adjusted HR 0.91; 95% CI, 0.73 to 1.14; $P = 0.41$). A greater number of positive lymph nodes and more extensive local tumour invasion were the strongest independent predictors of death. The proportional-hazards assumption was satisfied for the treatment effect (global Schoenfeld test, $P > 0.05$), so the reported hazard ratios are a valid summary across the follow-up period.

The advantage of the combination arm was also evident, and emerged earlier, for recurrence-free survival, indicating that the survival gain is driven predominantly by a reduction in disease recurrence rather than by effects after recurrence has occurred.

A Random Survival Forest yielded discrimination comparable to the Cox model (C-index 0.62 versus 0.66) and prioritised the same prognostic variables, supporting the adequacy and parsimony of the Cox model as the primary inferential summary.

8 Conclusions and clinical implications

In this randomised trial, adjuvant levamisole plus fluorouracil significantly delayed recurrence and prolonged overall survival in stage B and C colon cancer, reducing the rate of death by roughly one third relative to observation, whereas levamisole alone did not. The benefit was greatest among patients at the highest baseline risk (more than four positive nodes and greater local extent), who represent the clearest candidates for combination therapy. Methodologically, the agreement between a flexible machine-learning model and the conventional Cox model supports reporting the interpretable hazard ratios as the headline result.

9 Limitations

This is a single historical trial, so the specific drug regimen reflects the standard of its era rather than today's protocols; the value here is the method, not a current treatment recommendation. Some covariates have missing values (handled by complete-case analysis for the models). Survival analysis estimates associations adjusted for measured factors only; unmeasured confounding is controlled by randomisation for the treatment comparison but not for the observational prognostic factors.

Built in R (survival, survminer, randomForestSRC, gtsummary) from the colon cancer adjuvant-therapy trial, a real published RCT with openly available patient-level data. Author: Kingsley Amegah.